

cnn_tf_keras

August 1, 2021

1 Digit Recognition

1.1 Convolutional Neural Network implementation using TensorFlow - Keras

- Venkata A. Bharat Pothavajhala

```
[1]: import tensorflow as tf
     tf.__version__
```

```
[1]: '2.4.0'
```

```
[2]: import tensorflow.keras as keras
     print('TensorFlow version:',tf.__version__)
     print('Keras version:',keras.__version__)
```

TensorFlow version: 2.4.0

Keras version: 2.4.0



```
[3]: mnist = tf.keras.datasets.mnist.load_data(path='mnist.npz')
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
[25]: import matplotlib.pyplot as plt
     import matplotlib.image as mpimg
     import os
     %matplotlib inline
```

```
[26]: (train_X, train_Y), (test_X, test_Y) = mnist
```

```
[29]: for i in range(2):
     plt.imshow(train_X[i])
     plt.show()
```



```
[6]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense
```

```
[7]: model = Sequential()
inputShape = (28,28,1)
# The input layer accepts an image and applies a convolution that uses 32 6x6
↳ filters and a rectified linear unit activation function
model.add(Conv2D(32,
↳ kernel_size=(6,6),strides=(1,1),activation='relu',input_shape = inputShape))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(32, kernel_size=(6,6),strides=(1,1),activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
# model.add(Conv2D(32, kernel_size=(3,3),strides=(1,1),activation='relu'))
# model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.1))
model.add(Flatten())
model.add(Dense(10,activation='softmax'))
```

```
[8]: # With the layers defined, we can now compile the model for categorical
↳ (multi-class) classification
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 23, 23, 32)	1184
max_pooling2d (MaxPooling2D)	(None, 11, 11, 32)	0
conv2d_1 (Conv2D)	(None, 6, 6, 32)	36896
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 32)	0
dropout (Dropout)	(None, 3, 3, 32)	0
flatten (Flatten)	(None, 288)	0
dense (Dense)	(None, 10)	2890

Total params: 40,970
 Trainable params: 40,970
 Non-trainable params: 0

None

```
[9]: train_X.shape
train_X = train_X.reshape(-1,28,28,1)
train_X.shape
```

```
[9]: (60000, 28, 28, 1)
```

```
[10]: #import numpy as np
y = tf.keras.utils.to_categorical(train_Y, num_classes=10)
y
```

```
[10]: array([[0., 0., 0., ..., 0., 0., 0.],
          [1., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

```
[11]: # Train the model over 5 epochs using 30-image batches and using the validation
      ↪holdout dataset for validation
num_epochs = 5
batch_size = 30
history = model.fit(train_X,y,batch_size=batch_size,epochs = num_epochs)
```

```
Epoch 1/5
2000/2000 [=====] - 43s 21ms/step - loss: 1.7607 -
accuracy: 0.8177
Epoch 2/5
2000/2000 [=====] - 43s 21ms/step - loss: 0.1073 -
accuracy: 0.9689
Epoch 3/5
2000/2000 [=====] - 42s 21ms/step - loss: 0.0813 -
accuracy: 0.9756
Epoch 4/5
2000/2000 [=====] - 43s 21ms/step - loss: 0.0771 -
accuracy: 0.9770
Epoch 5/5
2000/2000 [=====] - 43s 21ms/step - loss: 0.0658 -
accuracy: 0.9813
```

```
[12]: model.save('nn_tfTrained.mdl')
```

```
INFO:tensorflow:Assets written to: nn_tfTrained.mdl/assets
```

```
[17]: #from google.colab import files

# with open('nn_tfTrained.mdl', 'w') as f:
#     f.write('some content')
```

```
files.download('/content/nn_tfTrained.mdl')
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.Javascript object>
```

```
[ ]: testY = tf.keras.utils.to_categorical(test_Y,num_classes = 10)
```

```
[ ]: test_X[1].shape[1]
```

```
[ ]: 28
```

```
[ ]: from tensorflow.keras import models
import numpy as np
from random import randint
import os
%matplotlib inline
```

```
[ ]: classnames = list(range(10))
classnames
```

```
[ ]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
[ ]: # Function to predict the class of an image
def predict_image(classifier, image):
    from tensorflow import convert_to_tensor
    # The model expects a batch of images as input, so we'll create an array of
    ↪ 1 image
    imgfeatures = img.reshape(1, img.shape[0], img.shape[1], 1)

    # We need to format the input to match the training data
    # The generator loaded the values as floating point numbers
    # and normalized the pixel values, so...
    imgfeatures = imgfeatures.astype('float32')
    #imgfeatures /= 255

    # Use the model to predict the image class
    class_probabilities = classifier.predict(imgfeatures)

    # Find the class predictions with the highest predicted probability
    index = int(np.argmax(class_probabilities, axis=1)[0])
    return index


# Create a random test image
img = test_X[4]
```

```
plt.axis('off')
plt.imshow(img)

# Use the classifier to predict the class
model1 = models.load_model('nn_tfTrained.mdl') # loads the saved model
class_idx = predict_image(model1, img)
print('index is ',class_idx)
print ('class is ',classnames[class_idx])
```

WARNING:tensorflow:5 out of the last 5 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7fd19c9961e0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has experimental_relax_shapes=True option that relaxes argument shapes that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

index is 4
class is 4



[]: